

Operating Systems

User Space Preemptive Scheduler

Giorgio Grisetti

`grisetti@diag.uniroma1.it`

Department of Computer Control and Management Engineering
Sapienza University of Rome

In this episode

- We will construct on the structure laid out in the previous lesson, a user space preemptive scheduler
- We need
 - a surrogate of the timer interrupt (we will use signals)
 - contexts for coroutines to store the possible "stacks" of our system
 - N contexts (one per process)
 - 1 context to handle the OS trap
 - 1 context to handle the timer interrupt
- Our system is purely user space.
 - We cannot implement the fork directly since it would involve copying "the memory" and preserving the links. Too complicated without access to the page table.
 - We will replace the fork() with a spawn(<function pointer>) similar in behavior to pthread create
 - Our spawn preserves the parent-child relationship of fork

Contexts

- 1 context per process (together with its stack), we put it in the PCB
- 1 context for the "interrupt"
- 1 context for the OS trap
- 1 context for the main, used on shutdown

```
typedef struct PCB{
    ListItem list; // MUST BE THE FIRST!!!
    int pid;
    int return_value; // ret value for the parent
    ProcessStatus status;
    int signals;
    int signals_mask;
    ListHead descriptors;
    struct PCB* parent;
    ListHead children;

    ucontext_t cpu_state; // the context
    char stack[STACK_SIZE]; // the context stack

    int syscall_num;
    long int syscall_args[DSOS_MAX_SYSCALLS_ARGS];
    int syscall_retvalue;
} PCB;

// in disastrOS.c, global variables

ucontext_t interrupt_context;
ucontext_t trap_context;
ucontext_t main_context;
```

Interrupts

We will mimic the interrupts with SIGALARM.

- SIGALARM is a signal that can be programmed to be sent periodically to our process
- The signal handler for sigalarm, will switch to the interrupt context,
- The interrupt is a function that will call the scheduler and "jump" transfer control to the next process, whose context is stored in `running->cpu_state`

```
//set up the signal action
void setupSignals(void) {
    struct sigaction act;

    // timerHandler is the function called
    // when signal is received!
    act.sa_sigaction = timerHandler;

    // restart the signal handler,
    // and take the handler from the
    // sa_sigaction field
    act.sa_flags = SA_RESTART | SA_SIGINFO;

    // handle only sigalarm
    sigemptyset(&act.sa_mask);
    sigemptyset(&signal_set);
    sigaddset(&signal_set, SIGALRM);

    // install the handler
    if(sigaction(SIGALRM, &act, NULL) != 0) {
        perror("Signal handler");
    }

    // start a system timer that will raise
    // a sigalarm each INTERVAL ms
    struct itimerval it;
    it.it_interval.tv_sec = 0;
    it.it_interval.tv_usec = INTERVAL * 1000;
    it.it_value = it.it_interval;
    if (setitimer(ITIMER_REAL, &it, NULL) )
        perror("setitimer");
}
```

Interrupts

We will mimic the interrupts with SIGALARM.

- SIGALARM is a signal that can be programmed to be sent periodically to our process
- The signal handler for sigalarm, will switch to the interrupt context,
- The interrupt is a function that will call the scheduler and "jump" transfer control to the next process, whose context is stored in `running->cpu_state`

```
void timerHandler(int j,
                  siginfo_t *si,
                  void *old_context) {
    // this saves the running context in the PCB
    swapcontext(&running->cpu_state,
               &interrupt_context);
}

void timerInterrupt(){

    ++disastrOS_time;
    printf("time: %d\n", disastrOS_time);

    // call the scheduler!!!
    internal_schedule();

    // this jumps to the pcb context
    // WITHOUT saving the state of the interrupt
    // next calls to the interrupt will restart
    // from the beginning of the function!
    setcontext(&running->cpu_state);
}
```

Trap

When invoking the syscall, we will

- pack the arguments in registers
- trap to the OS (swapping context). The swap operation will save the state in the process PCB

The trap will

- Decode the syscall, and call the appropriate routine in the syscall vector
- jump to the context of the running PCB

```
int disastrOS_syscall(int syscall_num, ...) {
    assert(running);
    va_list ap;
    if (syscall_num<0||syscall_num>DSOS_MAX_SYSCALLS)
        return DSOS_ESYSCALL_OUT_OF_RANGE;

    // pack the arguments of the
    // syscalls in the running pcb
    int nargs=syscall_numarg[syscall_num];
    va_start(ap,syscall_num);
    for (int i=0; i<nargs; ++i){
        running->syscall_args[i] = va_arg(ap,long int);
    }
    va_end(ap);
    running->syscall_num=syscall_num;

    // save the state of the CPU and
    // in PCB and jump to the trap
    swapcontext(&running->cpu_state, &trap_context);
    return running->syscall_retvalue;
}
```

Trap

When invoking the syscall, we will

- pack the arguments in registers
- trap to the OS (swapping context). the swap operation will save the state in the process PCB

The trap will

- Decode the syscall, and call the appropriate routine in the syscall vector
- jump to the context of the running PCB

```
void disastrOS_trap(){
    int syscall_num=running->syscall_num;

    if (syscall_num<0||syscall_num>DSOS_MAX_SYSCALLS){
        running->syscall_retvalue
        = DSOS_ESYSCALL_OUT_OF_RANGE;
        goto return_to_process;
    }
    SyscallFunctionType my_syscall
    =syscall_vector[syscall_num];
    if (! my_syscall) {
        running->syscall_retvalue
        =DSOS_ESYSCALL_NOT_IMPLEMENTED;
        goto return_to_process;
    }

    (*syscall_vector[syscall_num])();

return_to_process:
    if (running) {
        // trap to the process context saved
        // before the syscall
        setcontext(&running->cpu_state);
    } else {
        printf("no active processes\n");
        disastrOS_printStatus();
    }
}
```

Initializing Contexts

- We need to set the trap context so that it executes the `disastrOS_trap()`

- The trap context should "mask" the timer interrupts. This is done by setting the signal mask in the context

- We need to set the interrupt context so that it executes `timerInterrupt()`

```
void disastrOS_start(void (*f)(void*),
                    void* f_args, char* logfile){
    // .....

    // we will come back here on shutdown
    getcontext(&main_context);
    if (shutdown_now)
        exit(0);

    // setting system trap
    getcontext(&trap_context);
    trap_context.uc_stack.ss_sp = system_stack;
    trap_context.uc_stack.ss_size = STACK_SIZE;

    // we mask sigalarm when handing a trap
    sigemptyset(&trap_context.uc_sigmask);
    sigaddset(&trap_context.uc_sigmask, SIGALRM);
    trap_context.uc_stack.ss_flags = 0;
    trap_context.uc_link = &main_context;
    makecontext(&trap_context, disastrOS_trap, 0);

    // the interrupt and the system
    // share the same stack
    interrupt_context=trap_context;
    interrupt_context.uc_link = &main_context;
    sigemptyset(&interrupt_context.uc_sigmask);
    makecontext(&interrupt_context, timerInterrupt, 0);

    // .....
}
```


spawn

- We substitute the fork with a spawn routine that starts a thread.

- A thread has the following prototype

```
void f(void* arg)
```

- Spawn creates a new PCB and preserves the parent child-relation as done by fork. The new PCB is put in ready state.
- The context in the newly created PCB should accept timer interrupts.
- The stack is stored in the PCB

The syscall arguments are:

- args[0]: f_ptr
- args[1]: arg
- The newly created context will start the function pointed by f_ptr, with arguments arg

```
void internal_spawn(){
    static PCB* new_pcb;
    new_pcb=PCB_alloc();
    if (!new_pcb) {
        running->syscall_retvalue=DSOS_ESPAWN;
        return;
    }

    new_pcb->status=Ready;
    new_pcb->parent=running;
    PCBPtr* new_pcb_ptr=PCBPtr_alloc(new_pcb);
    assert(new_pcb_ptr);
    List_insert(&running->children,
                running->children.last,
                (ListItem*) new_pcb_ptr);
    List_insert(&ready_list,
                ready_list.last,
                (ListItem*) new_pcb);

    running->syscall_retvalue=new_pcb->pid;

    getcontext(&new_pcb->cpu_state);
    new_pcb->cpu_state.uc_stack.ss_sp=new_pcb->stack;
    new_pcb->cpu_state.uc_stack.ss_size=STACK_SIZE;
    new_pcb->cpu_state.uc_stack.ss_flags=0;
    sigemptyset(&new_pcb->cpu_state.uc_sigmask);
    new_pcb->cpu_state.uc_link = &main_context;
    void (*new_function)(void*)=
        (void(*) (void*))running->syscall_args[0];
    makecontext(&new_pcb->cpu_state,
                (void(*)())new_function,
                1,
                (void*)running->syscall_args[1]);
}
```

Run baby run

We just implemented an user space preemptive scheduler.

We can now run a bunch of threads

There will be a thread listening the keyboard and printing the status on the screen each time we press enter

There will be a bunch of threads executing dummy iterations. The number of iterations depends on the PID

Init

- will spawn the threads
- wait for the termination of all threads

When done it will call `disastrOS_shutdown()` to return to the main.

The main will just start the system, with the `disastrOS_start`.

```
void waitABit() {
    for (int i=0; i<1000000000; ++i);
}
// we need this to handle the sleep state
void sleeperFunction(void* args){
    printf("Hello, I am the sleeper,
           and I sleep %d\n",disastrOS_getpid());
    while(1) {
        getc(stdin);
        disastrOS_printStatus();
    }
}

void childFunction(void* args){
    printf("Hello, I am the child
           function %d\n",disastrOS_getpid());
    printf("I will iterate a bit,
           before terminating\n");
    for (int i=0; i<(disastrOS_getpid()+1); ++i){
        printf("PID: %d, iterate %d\n",
               disastrOS_getpid(), i);
        waitABit();
    }
    printf("PID: %d, terminating\n",
           disastrOS_getpid());
    disastrOS_exit(disastrOS_getpid()+1);
}
```

Run baby run

We just implemented an user space preemptive scheduler.

We can now run a bunch of threads

There will be a thread listening the keyboard and printing the status on the screen each time we press enter

There will be a bunch of threads executing dummy iterations. The number of iterations depends on the PID

Init

- will spawn the threads
- wait for the termination of all threads

When done it will call `disastrOS_shutdown()` to return to the main.

The main will just start the system, with the `disastrOS_start`.

```
void initFunction(void* args) {
    disastrOS_printStatus();
    printf("hello, I am init and I just started\n");
    disastrOS_spawn(sleeperFunction, 0);

    printf("I feel like to spawn 10 nice threads\n");
    int alive_children=0;
    for (int i=0; i<10; ++i) {
        disastrOS_spawn(childFunction, 0);
        alive_children++;
    }

    disastrOS_printStatus();

    printf("waiting for childs to terminate...\n");
    int retval;
    int pid;
    while(alive_children>0 &&
        (pid=disastrOS_wait(0, &retval))>=0){
        disastrOS_printStatus();
        printf("initFunction,
                child: %d terminated,
                retval:%d, alive: %d \n",
                pid, retval, alive_children);
        waitABit();
        --alive_children;
    }
    printf("shutdown!");
    disastrOS_shutdown();
}
```

Run baby run

We just implemented an user space preemptive scheduler.

We can now run a bunch of threads

There will be a thread listening the keyboard and printing the status on the screen each time we press enter

There will be a bunch of threads executing dummy iterations. The number of iterations depends on the PID

Init

- will spawn the threads
- wait for the termination of all threads

When done it will call `disastrOS_shutdown()` to return to the main.

The main will just start the system, with the `disastrOS_start`.

```
int main(int argc, char** argv){
    char* logfilename=0;
    if (argc>1) {
        logfilename=argv[1];
    }
    // we create the init process processes
    // the first is in the running variable
    // the others are in the ready queue
    printf("the function pointer is: %p",
        childFunction);
    // spawn an init process
    printf("start\n");
    disastrOS_start(initFunction, 0, logfilename);
    return 0;
}
```

Exercises

Modify the scheduler routine and the spawn function so that each thread can be given a priority.

The scheduler will pick first the threads with higher priority.

Once a thread is picked, its priority value is decreased each time the thread is evicted.

When the priority is 0, its value is restored to the initial priority.